

```
In [1]: 1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import plotly.express as px
6 import plotly.graph_objects as go
7 import plotly.io as pio
8 from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
9 from pandas.plotting import lag_plot, autocorrelation_plot
10 from statsmodels.tsa.seasonal import seasonal_decompose
11 from statsmodels.tsa.stattools import grangercausalitytests
12 from dateutil.parser import parse
13 from statsmodels.tsa.stattools import adfuller
14 from sklearn.preprocessing import MinMaxScaler, StandardScaler
15 from sklearn.metrics import mean_squared_error, mean_absolute_error
16 from keras.models import Sequential
17 from keras.layers import LSTM, Dense
18 from statsmodels.tsa.arima.model import ARIMA
19 import time
20
21 import scipy
22 import csv
23 import getpass
24 import pyodbc
25
26 import warnings
27 pio.renderers.default = 'notebook'
28 warnings.filterwarnings("ignore")
29 %matplotlib inline
```

2024-10-15 15:58:39.955528: I external/local\_tsl/tsl/cuda/cudart\_stub.cc:32] Could not find cuda drivers on your machine, GPU will not be used.

2024-10-15 15:58:40.079122: I external/local\_tsl/tsl/cuda/cudart\_stub.cc:32] Could not find cuda drivers on your machine, GPU will not be used.

2024-10-15 15:58:40.457359: I tensorflow/core/platform/cpu\_feature\_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

2024-10-15 15:58:42.072525: W tensorflow/compiler/tf2tensorrt/utils/py\_utils.cc:38] TF-TRT Warning: Could not find TensorRT

## • Connect into SQL-Database server

```
In [2]: 1 server = '192.168.122.56\Mohamed'
2 database = 'Sales_OLTP'
3 username = getpass.getpass("Please Enter your Login")
4 password = getpass.getpass("Please Enter your Password")
5 port = '1433'
6 Path = "/home/mohamed/Desktop/sales.csv"
7
8 connection_string = f'DRIVER={{ODBC Driver 17 for SQL Server}};SERVER={server},{port};D
9
10 try:
11     connection = pyodbc.connect(connection_string)
12
13     cursor = connection.cursor()
14
15     query = """SELECT DueDate,sum(TotalDue) as TotalDue
16                from Sales_OLTP.Sales.SalesOrderHeader
17                Group by DueDate
18                Order by DueDate
19            """
20
21     cursor.execute(query)
22
23     rows = cursor.fetchall()
24
25     csv_file_path = 'sales_data.csv'
26
27     with open(Path, mode='w', newline='', encoding='utf-8') as file:
28         writer = csv.writer(file)
29
30         writer.writerow([desc[0] for desc in cursor.description])
31
32         for row in rows:
33             writer.writerow(row)
34
35     print(f"Data has been exported successfully to {csv_file_path}")
36
37     cursor.close()
38     connection.close()
39
40 except Exception as e:
41     print(f"Error: {e}")
42
```

Please Enter your Login.....

Please Enter your Password.....

Data has been exported successfully to sales\_data.csv

## • Loading data¶

```
In [2]: 1 df = pd.read_csv("/home/mohamed/Desktop/sales.csv",index_col="DueDate")
2 df.index = pd.to_datetime(df.index)
```

```
In [3]: 1 df = df.tz_localize("UTC").tz_convert("Africa/Cairo")
```

- Description of data

```
In [4]: 1 print(df.info())
        2 df.describe()
```

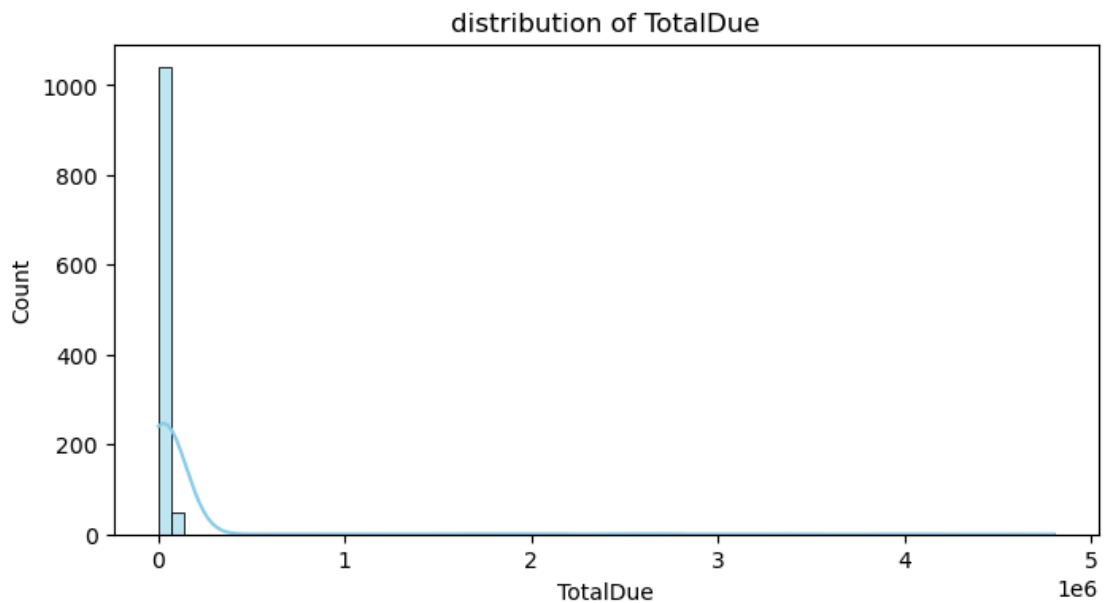
```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1124 entries, 2011-06-12 02:00:00+02:00 to 2014-07-12 02:00:00+02:00
Data columns (total 1 columns):
 #   Column      Non-Null Count  Dtype  
---  --
 0   TotalDue    1124 non-null   float64
dtypes: float64(1)
memory usage: 17.6 KB
None
```

Out[4]:

	TotalDue
count	1.124000e+03
mean	1.096235e+05
std	4.845919e+05
min	7.725036e+02
25%	1.559133e+04
50%	2.382283e+04
75%	4.171030e+04
max	4.800611e+06


- visualize time-series data

```
In [5]: 1 plt.figure(figsize=(8, 4))
        2 sns.histplot(df.TotalDue, bins=70, kde=True, color="skyblue", edgecolor="black");
        3 plt.title("distribution of TotalDue")
        4 plt.show()
```




```
In [6]: 1 fig = px.line(df, x=df.index, y=df.TotalDue, title='TotalDue over Time')
2
3 fig.show()
```

TotalDue over Time



```
In [7]: 1 x = df.index
2 y1 = df["TotalDue"].values
3
4 fig, ax = plt.subplots(1, 1, figsize=(16,5), dpi= 120)
5 plt.fill_between(x, y1=y1, y2=-y1, alpha=0.5, linewidth=2, color='seagreen')
6 plt.title('TotalDue', fontsize=16)
7 plt.hlines(y=0, xmin=np.min(x), xmax=np.max(x), linewidth=.5)
8 plt.show()
```



- from line chart above i can say this time-series is non stationary But i will verify this using the mathematical method (ADF).

## • Augmented Dickey-Fuller (ADF) test:

- First, I will check if the series is stationary using the Augmented Dickey Fuller test (ADF Test), from the statsmodels package. The reason being is that we need differencing only if the series is non-stationary. Else, no differencing is needed, that is,  $d=0$ .
- The null hypothesis ( $H_0$ ) of the ADF test is that the time series is non-stationary. So, if the p-value of the test is less than the significance level (0.05) then we reject the null hypothesis and infer that the time series is indeed stationary.


- So, in our case, if P Value > 0.05 we go ahead with finding the order of differencing.

```
In [8]: 1 result = adfuller(df)
2
3 print('ADF Statistic:', result[0])
4 print('p-value:', result[1])
5 print('Critical Values:')
6 for key, value in result[4].items():
7     print(f'\t\t{key}: {value}')
8
```

```
ADF Statistic: -34.359333112713145
p-value: 0.0
Critical Values:
1%: -3.4361864296062166
5%: -2.864117116658563
10%: -2.5681421294173714
```


- from ADF-test above ADF statistic is more negative than the critical values at the 1%, 5%, or 10% levels and p-value = 0.0 is greater than 0.05, so We Must reject the null hypothesis and conclude that the time series is stationary.
  - i can say Differencing equal 0 (d=0)
- I will find out the required number of AR terms by inspecting the Partial Autocorrelation (PACF) plot.

```
In [9]: 1 fig, axes = plt.subplots(figsize=(10,3), dpi= 100)
2 plot_pacf(df, lags=50, ax=axes)
3 plt.show()
```



- We can see that the PACF lag (P) from (30) TO (33) is quite significant since it is well above the significance line.
- we will look at the ACF plot for the number of MA terms. An MA term is technically, the error of the lagged forecast.

```
In [10]: 1 fig, axes = plt.subplots(figsize=(10,3), dpi= 100)
2 plot_acf(df, lags=50, ax=axes)
3 plt.show()
```



- We can see that the ACF lag (q) 1
- **Now, we have determined the values of p, d and q. We have everything needed to fit the ARIMA model. We will use the ARIMA() implementation in the statsmodels package.**
- hyperparameters
  - p -> (30-33)
  - q -> (30-31)
  - d -> 0

## • Training the model

```
In [11]: 1 lags = np.arange(30,34,1)
2 MAs = np.arange(30,32,1)
3 best_model = None
4 best_aic = float('inf')
5
6 for p in lags:
7     for q in MAs:
8         try:
9             model = ARIMA(df['TotalDue'], order=(p, 0, q))
10            model_fit = model.fit()
11
12            print(f"Fitted ARIMA({p},0,{q}) - AIC: {model_fit.aic}")
13
14            if model_fit.aic < best_aic:
15                best_aic = model_fit.aic
16                best_model = model_fit
17
18            except Exception as e:
19                print(f"Failed to fit ARIMA({p},0,{q}): {e}")
20
21 if best_model:
22     print("\nBest Model Summary:")
23     print(best_model.summary())
24 else:
25     print("No valid model found.")
26
```

Fitted ARIMA(30,0,30) - AIC: 32166.897912468557  
 Fitted ARIMA(30,0,31) - AIC: 32089.213183022435  
 Fitted ARIMA(31,0,30) - AIC: 32027.200164254522  
 Fitted ARIMA(31,0,31) - AIC: 32050.843698195888  
 Fitted ARIMA(32,0,30) - AIC: 32022.62102025915  
 Fitted ARIMA(32,0,31) - AIC: 32042.500686776788  
 Fitted ARIMA(33,0,30) - AIC: 32013.67674092082  
 Fitted ARIMA(33,0,31) - AIC: 32008.658113045793

Best Model Summary:

#### SARIMAX Results

```

=====
Dep. Variable:          TotalDue      No. Observations:      1124
Model:                ARIMA(33, 0, 31)  Log Likelihood        -15938.329
Date:                 Tue, 15 Oct 2024  AIC                  32008.658
Time:                 16:12:17         BIC                  32340.285
Sample:               0                HQIC                 32133.982
                             - 1124
Covariance Type:      opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	1.096e+05	3.48e-07	3.15e+11	0.000	1.1e+05	1.1e+05
ar.L1	-1.0092	0.192	-5.258	0.000	-1.385	-0.633
ar.L2	-0.4291	0.095	-4.509	0.000	-0.616	-0.243
ar.L3	-0.1022	0.078	-1.307	0.191	-0.256	0.051
ar.L4	-0.0173	0.084	-0.206	0.837	-0.182	0.147
ar.L5	-0.0313	0.093	-0.336	0.737	-0.214	0.151
ar.L6	0.0012	0.107	0.011	0.991	-0.209	0.212
ar.L7	-0.0317	0.110	-0.289	0.773	-0.247	0.184
ar.L8	-0.0025	0.125	-0.020	0.984	-0.248	0.243
ar.L9	-0.0169	0.117	-0.145	0.885	-0.245	0.211
ar.L10	-0.0127	0.157	-0.081	0.936	-0.321	0.295
ar.L11	-0.0126	0.134	-0.095	0.925	-0.275	0.250
ar.L12	-0.0155	0.172	-0.090	0.928	-0.353	0.322
ar.L13	-0.0169	0.155	-0.109	0.913	-0.320	0.286
ar.L14	-0.0072	0.162	-0.044	0.965	-0.326	0.311
ar.L15	-0.0194	0.164	-0.118	0.906	-0.341	0.302
ar.L16	-0.0067	0.165	-0.041	0.967	-0.330	0.316
ar.L17	-0.0167	0.165	-0.101	0.920	-0.340	0.307
ar.L18	-0.0091	0.161	-0.056	0.955	-0.324	0.306
ar.L19	-0.0132	0.134	-0.098	0.922	-0.276	0.250
ar.L20	-0.0075	0.125	-0.060	0.952	-0.252	0.237
ar.L21	-0.0170	0.115	-0.148	0.882	-0.242	0.208
ar.L22	-0.0009	0.099	-0.010	0.992	-0.195	0.193
ar.L23	-0.0264	0.090	-0.292	0.770	-0.204	0.151
ar.L24	0.0104	0.089	0.117	0.907	-0.164	0.185
ar.L25	-0.0390	0.074	-0.525	0.600	-0.185	0.107
ar.L26	0.0279	0.067	0.419	0.675	-0.103	0.158
ar.L27	-0.0545	0.068	-0.800	0.424	-0.188	0.079
ar.L28	0.0374	0.067	0.556	0.578	-0.095	0.169
ar.L29	-0.0216	0.071	-0.303	0.762	-0.161	0.118
ar.L30	0.4718	0.057	8.227	0.000	0.359	0.584
ar.L31	1.1177	0.131	8.543	0.000	0.861	1.374
ar.L32	0.6231	0.105	5.948	0.000	0.418	0.828
ar.L33	0.2301	0.063	3.656	0.000	0.107	0.354
ma.L1	0.6597	0.192	3.428	0.001	0.283	1.037
ma.L2	0.0516	0.100	0.514	0.607	-0.145	0.248
ma.L3	-0.0479	0.073	-0.657	0.511	-0.191	0.095
ma.L4	0.0293	0.106	0.276	0.782	-0.178	0.237
ma.L5	0.0821	0.128	0.642	0.521	-0.169	0.333
ma.L6	0.0262	0.158	0.166	0.868	-0.284	0.336
ma.L7	0.0670	0.170	0.395	0.693	-0.265	0.399
ma.L8	0.0095	0.197	0.049	0.961	-0.376	0.395
ma.L9	0.0268	0.205	0.131	0.896	-0.374	0.428
ma.L10	0.0203	0.254	0.080	0.936	-0.477	0.518
ma.L11	0.0181	0.246	0.074	0.941	-0.464	0.501
ma.L12	0.0290	0.304	0.096	0.924	-0.566	0.624
ma.L13	0.0210	0.277	0.076	0.939	-0.521	0.563
ma.L14	0.0119	0.305	0.039	0.969	-0.586	0.610
ma.L15	0.0300	0.295	0.102	0.919	-0.549	0.609
ma.L16	0.0057	0.301	0.019	0.985	-0.584	0.595
ma.L17	0.0255	0.302	0.084	0.933	-0.567	0.618
ma.L18	0.0100	0.276	0.036	0.971	-0.531	0.551
ma.L19	0.0103	0.246	0.042	0.967	-0.472	0.492
ma.L20	0.0115	0.224	0.051	0.959	-0.427	0.450
ma.L21	0.0137	0.209	0.065	0.948	-0.395	0.422



ma.L22	-0.0026	0.180	-0.014	0.988	-0.356	0.351
ma.L23	0.0364	0.161	0.225	0.822	-0.280	0.353
ma.L24	-0.0337	0.136	-0.248	0.804	-0.300	0.232
ma.L25	0.0461	0.111	0.415	0.678	-0.171	0.264
ma.L26	-0.0689	0.091	-0.757	0.449	-0.247	0.110
ma.L27	0.0123	0.101	0.122	0.903	-0.186	0.211
ma.L28	-0.1095	0.073	-1.493	0.135	-0.253	0.034
ma.L29	-0.1031	0.080	-1.285	0.199	-0.261	0.054
ma.L30	-0.5782	0.050	-11.649	0.000	-0.675	-0.481
ma.L31	-0.6741	0.142	-4.738	0.000	-0.953	-0.395
sigma2	1.221e+11	1.83e-11	6.68e+21	0.000	1.22e+11	1.22e+11

---

Ljung-Box (L1) (Q):	0.14	Jarque-Bera (JB):	30784.57
Prob(Q):	0.71	Prob(JB):	0.00
Heteroskedasticity (H):	2.05	Skew:	2.91
Prob(H) (two-sided):	0.00	Kurtosis:	27.97

---

## Warnings:



[1] Covariance matrix calculated using the outer product of gradients (complex-step).  
 [2] Covariance matrix is singular or near-singular, with condition number 1.66e+37. Standard errors may be unstable.

- from this training will choose the model with parameters (33,0,31)

- plot the Residuals after training this model

```
In [12]: 1 residuals = pd.DataFrame(best_model.resid, columns=['Residuals'])
2
3 # Plot residuals
4 fig1 = px.line(residuals,
5               y='Residuals',
6               title='Residuals Over Time',
7               labels={'index': 'Observation', 'Residuals': 'Residual Value'})
8
9 fig2 = residuals.plot(kind='kde', title='Density')
10 fig1.show()
11 # fig2.show()
```

Residuals Over Time




```
In [13]: 1 # model = ARIMA(df['TotalDue'], order=(33, 0, 31))
2 # model_fit = model.fit()
```

### • prediction

```
In [14]: 1 pred = best_model.get_prediction(start=0, end=len(df)-1)
2 df['Forecast'] = pred.predicted_mean
3
4 df_plot = df[['TotalDue', 'Forecast']].copy()
5 df_plot['Date'] = df_plot.index
6 df_plot = df_plot.melt(id_vars='Date', var_name='Type', value_name='Value')
7
8 fig = px.line(df_plot, x='Date', y='Value', color='Type',
9               title='ARIMA Model - Forecast vs Actual Data',
10              labels={'Date': 'Date', 'Value': 'TotalDue'})
11
12 fig.show()
```

ARIMA Model - Forecast vs Actual Data



### • split data and train the model again

```
In [15]: 1 split = int(0.8 * len(x))
2 train, test = df['TotalDue'][:split], df['TotalDue'][split:]
```

```
In [16]: 1 model = ARIMA(train, order=(33, 0, 31))
2 fitted = model.fit()
```

- forecasting and validating th model


```
In [17]: 1 forecast_result = fitted.get_forecast(steps=119, alpha=0.05)
          2
          3 fc = forecast_result.predicted_mean
          4 conf = forecast_result.conf_int()
          5
          6 fc_series = pd.Series(fc.values, index=test.index[:len(fc)])
          7 lower_series = pd.Series(conf.iloc[:, 0].values, index=test.index[:len(fc)])
          8 upper_series = pd.Series(conf.iloc[:, 1].values, index=test.index[:len(fc)])
          9
```

```

In [18]: 1 # Plot
2 fig = go.Figure()
3
4 fig.add_trace(go.Scatter(x=train.index, y=train, mode='lines', name='Training'))
5
6 fig.add_trace(go.Scatter(x=test.index, y=test, mode='lines', name='Actual', line=dict(c
7
8 fig.add_trace(go.Scatter(x=fc_series.index, y=fc_series, mode='lines', name='Forecast',
9
10 fig.add_trace(go.Scatter(
11     x=fc_series.index, y=upper_series, mode='lines',
12     line=dict(color='gray', width=0), showlegend=False))
13
14 fig.add_trace(go.Scatter(
15     x=fc_series.index, y=lower_series, mode='lines',
16     line=dict(color='gray', width=0), showlegend=False,
17     fill='tonexty', fillcolor='rgba(128, 128, 128, 0.3)'))
18
19 fig.update_layout(
20     title='Forecast vs Actuals',
21     xaxis_title='Date',
22     yaxis_title='Total Due',
23     legend=dict(x=0.01, y=0.99, bgcolor='rgba(255,255,255,0.7)', bordercolor='black', b
24     width=1000, height=600
25 )
26
27 fig.show()
28

```

Forecast vs Actuals



```

In [19]: 1 aligned_test = test.values[:len(fc)]
          2
          3
          4 def forecast_accuracy(forecast, actual):
          5     mpe = np.mean((forecast - actual) / actual)
          6     rmse = np.sqrt(np.mean((forecast - actual) ** 2))
          7     minmax = 1 - np.mean(np.min([forecast, actual], axis=0) / np.max([forecast, actual]
          8
          9     return {'mpe': mpe, 'rmse': rmse, 'minmax': minmax}
          10
          11 accuracy_metrics = forecast_accuracy(fc, aligned_test)
          12 accuracy_metrics
          13

```

```

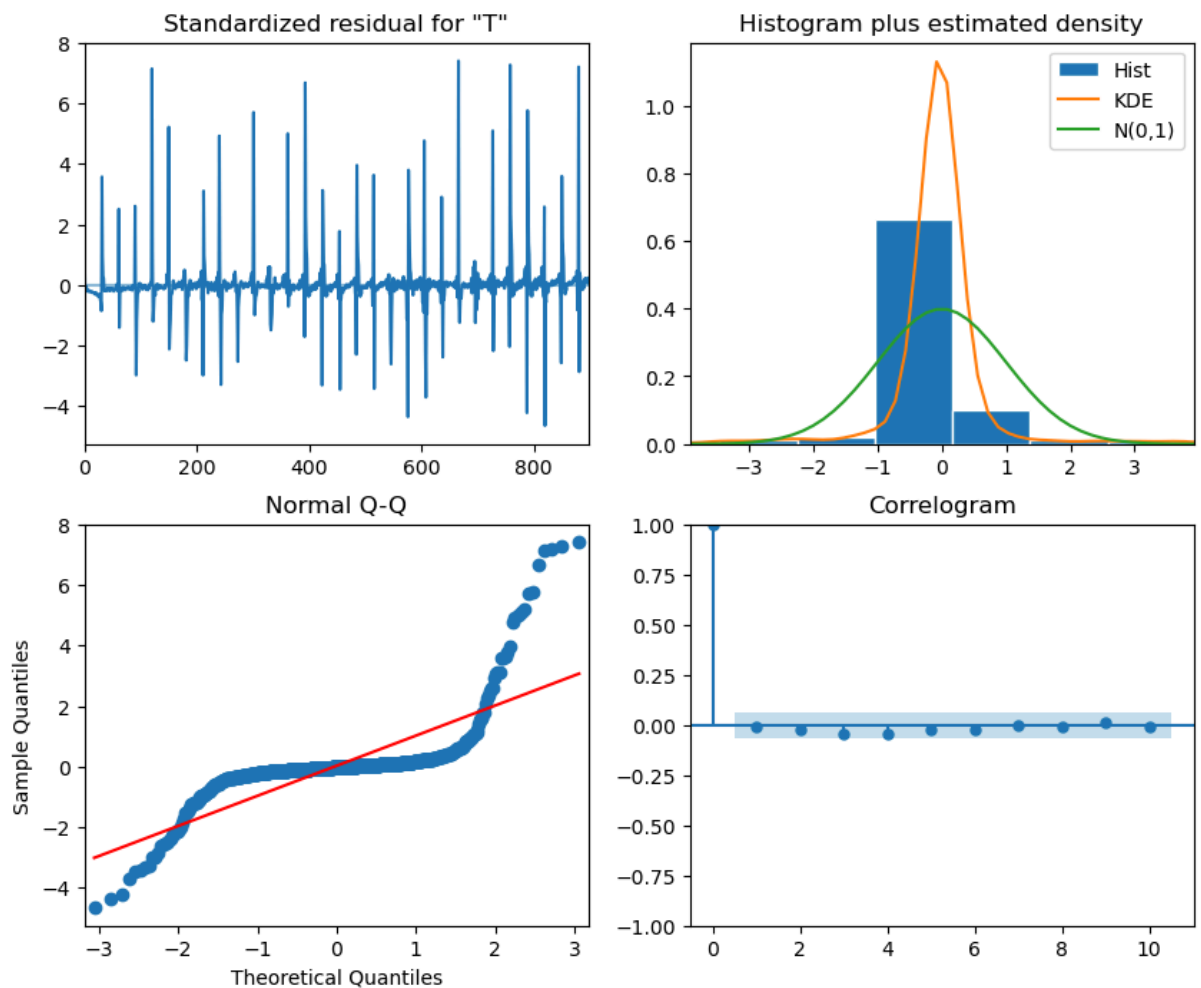
Out[19]: {'mpe': 0.728288216527285,
          'rmse': 444713.4826793588,
          'minmax': 1.4927666289249102}

```

```

In [20]: 1 fitted.plot_diagnostics(figsize=(10,8))
          2 plt.show()

```




- Standardized residual: The residual errors seem to fluctuate around a mean of zero and have a uniform variance.
- Histogram: The density plot suggest normal distribution with mean slightly shifted towards right.
- Theoretical Quantiles: Mostly the dots fall perfectly in line with the red line.
- Correlogram: The Correlogram, (or ACF plot) shows the residual errors are not autocorrelated.
- Overall, the model seems to be a good fit. So, let's use it to forecast but this need deseasonalized

```

In [21]: 1 steps = 150
2 forecast_result = fitted.get_forecast(steps=steps)
3
4 forecast = forecast_result.predicted_mean
5 conf_int = forecast_result.conf_int(alpha=0.05)
6
7 forecast_index = pd.date_range(start=test.index[-1], periods=steps + 1, freq='D')[1:]
8 forecast_series = pd.Series(forecast.values, index=forecast_index)
9
10 lower_series = pd.Series(conf_int.iloc[:, 0].values, index=forecast_index)
11 upper_series = pd.Series(conf_int.iloc[:, 1].values, index=forecast_index)
12
13
14 fig = go.Figure()
15
16 fig.add_trace(go.Scatter(x=train.index, y=train, mode='lines', name='Training', line=dict(
17 fig.add_trace(go.Scatter(x=test.index, y=test, mode='lines', name='test', line=dict(col
18
19 fig.add_trace(go.Scatter(x=forecast_series.index, y=forecast_series, mode='lines',
20 name='Forecast', line=dict(color='fuchsia'))))
21
22 fig.add_trace(go.Scatter(x=forecast_series.index, y=lower_series, mode='lines',
23 name='Lower Bound', line=dict(width=0), fill=None))
24 fig.add_trace(go.Scatter(x=forecast_series.index, y=upper_series, mode='lines',
25 name='Upper Bound', line=dict(width=0), fill='tonexty',
26 fillcolor='rgba(128, 128, 128, 0.2)', showlegend=False))
27
28 fig.update_layout(
29 title='Forecast vs Actuals with Future Predictions',
30 xaxis_title='Date',
31 yaxis_title='Total Due',
32 width=900, height=500,
33 legend=dict(yanchor='top', y=0.99, xanchor='left', x=0.01)
34 )
35
36 fig.show()
37

```

Forecast vs Actuals with Future Predictions




- From the analysis process of the above model, we notice that the model is good, but there are some problems in this model. The residuals are very large, and the model fits the data well. It is clear to us that this is due to the seasonality present in the data.

- deseasonalized


```
In [22]: 1 multiplicative_decomposition = seasonal_decompose(df['TotalDue'], model='multiplicative
2
3 additive_decomposition = seasonal_decompose(df['TotalDue'], model='additive', period=30
4
5 plt.rcParams.update({'figure.figsize': (16,12)})
6 multiplicative_decomposition.plot().suptitle('Multiplicative Decomposition', fontsize=1
7 plt.tight_layout(rect=[0, 0.03, 1, 0.95])
8
9 additive_decomposition.plot().suptitle('Additive Decomposition', fontsize=16)
10 plt.tight_layout(rect=[0, 0.03, 1, 0.95])
11
12 plt.show()
```

Multiplicative Decomposition



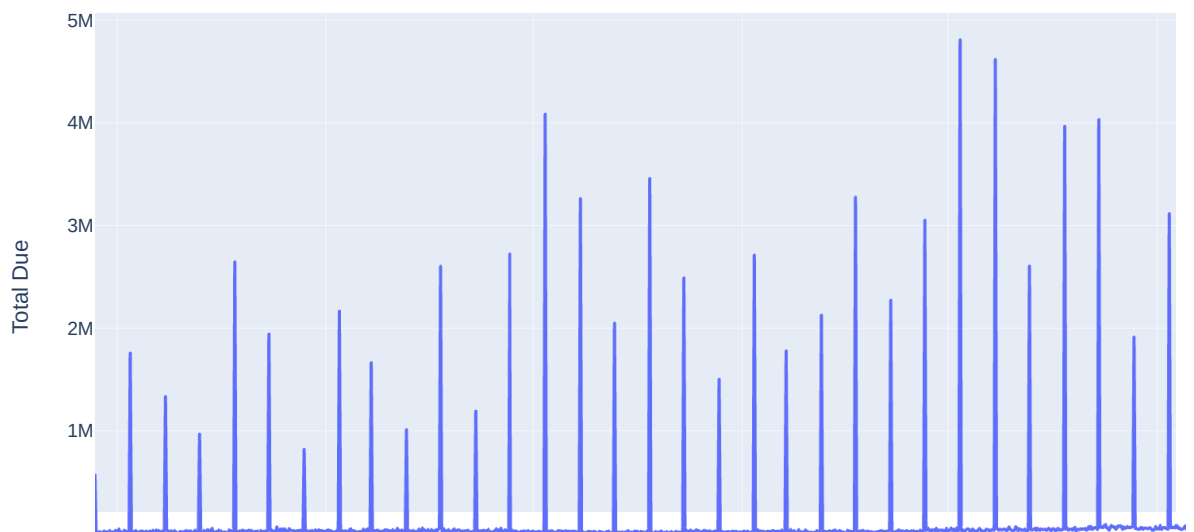


## Additive Decomposition



```
In [23]: 1 deseasonalized = (df['TotalDue'].dropna() / multiplicative_decomposition.seasonal).dropna()
2
3
4 # Plot
5 fig = px.line(deseasonalized,
6               x=df.index,
7               y=df.TotalDue,
8               title='deseasonalized TotalDue over Time',
9               labels={'x': 'DueDate', 'y': 'Total Due'})
10 fig.show()
```

deseasonalized TotalDue over Time




- check again stationary Augmented Dickey-Fuller (ADF) test and ACF and PACF

```
In [24]: 1 result = adfuller(deseasonalized)
2
3 print('ADF Statistic:', result[0])
4 print('p-value:', result[1])
5 print('Critical Values:')
6 for key, value in result[4].items():
7     print(f'\t\t{key}: {value}')
8
```

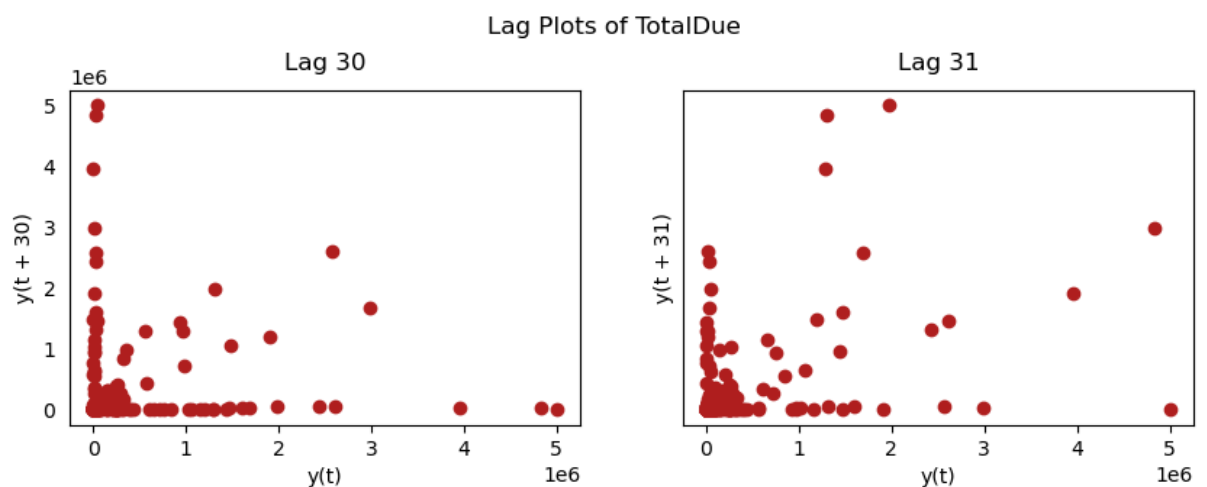
ADF Statistic: -34.059840588918846  
 p-value: 0.0  
 Critical Values:  
     1%: -3.4361864296062166  
     5%: -2.864117116658563  
    10%: -2.5681421294173714

```
In [25]: 1 fig, axes = plt.subplots(1,2,figsize=(16,3), dpi= 100)
2 plot_acf(deseasonalized.tolist(), lags=50, ax=axes[0])
3 plot_pacf(deseasonalized.tolist(), lags=50, ax=axes[1])
4 plt.show()
```



- plot lags

```
In [26]: 1 plt.rcParams.update({'ytick.left' : False, 'axes.titlepad':10})
2
3 fig, axes = plt.subplots(1, 2, figsize=(10,3), sharex=True, sharey=True, dpi=100)
4 for i, ax in enumerate(axes.flatten()[2]):
5     lag_plot(deseasonalized, lag=i+30, ax=ax, c='firebrick')
6     ax.set_title('Lag ' + str(i+30))
7
8 fig.suptitle('Lag Plots of TotalDue', y=1.05)
9 plt.show()
```



```
In [27]: 1 lags = np.arange(30,32,1)
2 MAs = np.arange(30,33,1)
3 best_model = None
4 best_aic = float('inf')
5
6 for p in lags:
7     for q in MAs:
8         try:
9             model = ARIMA(deseasonalized, order=(p, 0, q))
10            model_fit = model.fit()
11
12            print(f"Fitted ARIMA({p},0,{q}) - AIC: {model_fit.aic}")
13
14            if model_fit.aic < best_aic:
15                best_aic = model_fit.aic
16                best_model = model_fit
17
18            except Exception as e:
19                print(f"Failed to fit ARIMA({p},0,{q}): {e}")
20
21 if best_model:
22     print("\nBest Model Summary:")
23     print(best_model.summary())
24 else:
25     print("No valid model found.")
```

Fitted ARIMA(30,0,30) - AIC: 31529.953548983427  
 Fitted ARIMA(30,0,31) - AIC: 31420.082259036422  
 Fitted ARIMA(30,0,32) - AIC: 31398.343108113397  
 Fitted ARIMA(31,0,30) - AIC: 31296.16308961033  
 Fitted ARIMA(31,0,31) - AIC: 31309.099728508554  
 Fitted ARIMA(31,0,32) - AIC: 31309.490256154855

Best Model Summary:

#### SARIMAX Results

```
=====
Dep. Variable:          y          No. Observations:      1124
Model:                ARIMA(31, 0, 30)  Log Likelihood    -15585.082
Date:                 Tue, 15 Oct 2024  AIC                31296.163
Time:                  16:22:38         BIC                31612.716
Sample:                0              HQIC               31415.791
                             - 1124
Covariance Type:      opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	1.117e+05	3.08e-08	3.63e+12	0.000	1.12e+05	1.12e+05
ar.L1	-0.0784	0.039	-2.035	0.042	-0.154	-0.003
ar.L2	-0.0283	0.041	-0.691	0.490	-0.109	0.052
ar.L3	0.0020	0.050	0.040	0.968	-0.097	0.101
ar.L4	-0.0117	0.049	-0.239	0.811	-0.107	0.084
ar.L5	-0.0017	0.070	-0.025	0.980	-0.138	0.135
ar.L6	0.0085	0.067	0.125	0.900	-0.124	0.141
ar.L7	-0.0154	0.087	-0.178	0.858	-0.185	0.154
ar.L8	0.0156	0.085	0.184	0.854	-0.151	0.182
ar.L9	-0.0201	0.084	-0.239	0.811	-0.185	0.145
ar.L10	0.0227	0.080	0.286	0.775	-0.133	0.179
ar.L11	-0.0249	0.078	-0.320	0.749	-0.178	0.128
ar.L12	0.0129	0.092	0.140	0.889	-0.167	0.193
ar.L13	-0.0033	0.071	-0.046	0.963	-0.141	0.135
ar.L14	-0.0053	0.075	-0.070	0.944	-0.153	0.142
ar.L15	-0.0070	0.085	-0.082	0.935	-0.173	0.159
ar.L16	0.0023	0.082	0.028	0.978	-0.158	0.162
ar.L17	0.0088	0.059	0.148	0.882	-0.107	0.125
ar.L18	-0.0349	0.064	-0.546	0.585	-0.160	0.090
ar.L19	0.0243	0.069	0.352	0.725	-0.111	0.160
ar.L20	-0.0318	0.080	-0.399	0.690	-0.188	0.124
ar.L21	0.0185	0.095	0.193	0.847	-0.169	0.206
ar.L22	-0.0358	0.095	-0.378	0.706	-0.221	0.150
ar.L23	0.0180	0.104	0.173	0.862	-0.186	0.222
ar.L24	-0.0260	0.104	-0.251	0.801	-0.229	0.177
ar.L25	-0.0079	0.098	-0.081	0.935	-0.200	0.184
ar.L26	0.0188	0.087	0.215	0.830	-0.153	0.190
ar.L27	-0.0489	0.071	-0.693	0.488	-0.187	0.089
ar.L28	0.0569	0.049	1.162	0.245	-0.039	0.153
ar.L29	-0.1245	0.041	-3.074	0.002	-0.204	-0.045
ar.L30	0.6403	0.023	27.382	0.000	0.595	0.686
ar.L31	0.5039	0.022	23.003	0.000	0.461	0.547
ma.L1	-0.0269	0.046	-0.580	0.562	-0.118	0.064
ma.L2	0.0701	0.048	1.462	0.144	-0.024	0.164
ma.L3	-0.0252	0.069	-0.367	0.714	-0.160	0.109
ma.L4	0.0579	0.073	0.795	0.427	-0.085	0.201
ma.L5	0.0310	0.087	0.357	0.721	-0.139	0.202
ma.L6	0.0187	0.088	0.213	0.831	-0.153	0.190
ma.L7	0.0374	0.093	0.401	0.688	-0.145	0.220
ma.L8	-0.0126	0.091	-0.139	0.889	-0.190	0.165
ma.L9	0.0399	0.087	0.460	0.646	-0.130	0.210
ma.L10	-0.0051	0.093	-0.055	0.956	-0.187	0.177
ma.L11	0.0544	0.092	0.588	0.556	-0.127	0.236
ma.L12	-0.0105	0.088	-0.119	0.905	-0.182	0.161
ma.L13	0.0278	0.088	0.316	0.752	-0.145	0.200
ma.L14	0.0203	0.103	0.196	0.845	-0.182	0.223
ma.L15	0.0108	0.104	0.104	0.917	-0.194	0.215
ma.L16	0.0080	0.074	0.107	0.914	-0.137	0.153
ma.L17	0.0023	0.079	0.029	0.977	-0.152	0.157
ma.L18	0.0511	0.090	0.570	0.569	-0.125	0.227
ma.L19	-0.0023	0.086	-0.027	0.979	-0.170	0.165
ma.L20	0.0448	0.095	0.469	0.639	-0.142	0.232
ma.L21	-0.0105	0.098	-0.107	0.915	-0.203	0.182
ma.L22	0.0288	0.095	0.303	0.762	-0.158	0.216
ma.L23	0.0060	0.107	0.056	0.955	-0.203	0.215
ma.L24	0.0433	0.095	0.457	0.648	-0.142	0.229
ma.L25	0.0378	0.090	0.419	0.675	-0.139	0.215

ma.L26	-0.0173	0.096	-0.181	0.856	-0.205	0.170
ma.L27	0.0418	0.077	0.542	0.588	-0.109	0.193
ma.L28	-0.0210	0.059	-0.358	0.721	-0.136	0.094
ma.L29	0.1462	0.051	2.848	0.004	0.046	0.247
ma.L30	-0.8555	0.036	-23.759	0.000	-0.926	-0.785
sigma2	7.378e+10	2.96e-12	2.5e+22	0.000	7.38e+10	7.38e+10

---

Ljung-Box (L1) (Q):	0.05	Jarque-Bera (JB):	451753.09
Prob(Q):	0.83	Prob(JB):	0.00
Heteroskedasticity (H):	8.15	Skew:	6.52
Prob(H) (two-sided):	0.00	Kurtosis:	100.34

---


## Warnings:


[1] Covariance matrix calculated using the outer product of gradients (complex-step).  
 [2] Covariance matrix is singular or near-singular, with condition number 5.33e+37. Standard errors may be unstable.

- the best parameter for this model (31,0,30)

```
In [28]: 1 residuals = pd.DataFrame(best_model.resid, columns=['Residuals'])
2
3 # Plot residuals
4 fig1 = px.line(residuals,
5               y='Residuals',
6               title='Residuals Over Time',
7               labels={'index': 'Observation', 'Residuals': 'Residual Value'})
8
9 fig2 = plt.figure(figsize=(8, 6))
10 ax = residuals.plot(kind='kde', title='Density', ax=fig2.add_subplot(111))
11
12 fig1.show()
13 plt.show()
```

Residuals Over Time





- split deseasonalized data and train the model again

```
In [29]: 1 split = int(0.8 * len(x))  
2 train, test = deseasonalized.to_frame()[0][:split], deseasonalized.to_frame()[0][split:
```

```
In [30]: 1 model = ARIMA(train, order=(31, 0, 30))  
2 fitted = model.fit()
```


- Forecasting validation

```
In [31]: 1 forecast_result = fitted.get_forecast(steps=119, alpha=0.05)  
2  
3 fc = forecast_result.predicted_mean  
4 conf = forecast_result.conf_int()  
5  
6 fc_series = pd.Series(fc.values, index=test.index[:len(fc)])  
7 lower_series = pd.Series(conf.iloc[:, 0].values, index=test.index[:len(fc)])  
8 upper_series = pd.Series(conf.iloc[:, 1].values, index=test.index[:len(fc)])  
9
```

- validation test set

```
In [32]: 1 # Plot
2 fig = go.Figure()
3
4 fig.add_trace(go.Scatter(x=train.index, y=train, mode='lines', name='Training'))
5
6 fig.add_trace(go.Scatter(x=test.index, y=test, mode='lines', name='Actual', line=dict(c
7
8 fig.add_trace(go.Scatter(x=fc_series.index, y=fc_series, mode='lines', name='Forecast',
9
10 fig.add_trace(go.Scatter(
11     x=fc_series.index, y=upper_series, mode='lines',
12     line=dict(color='gray', width=0), showlegend=False))
13
14 fig.add_trace(go.Scatter(
15     x=fc_series.index, y=lower_series, mode='lines',
16     line=dict(color='gray', width=0), showlegend=False,
17     fill='tonexty', fillcolor='rgba(128, 128, 128, 0.3)'))
18
19 fig.update_layout(
20     title='Forecast vs Actuals',
21     xaxis_title='Date',
22     yaxis_title='Total Due',
23     legend=dict(x=0.01, y=0.99, bgcolor='rgba(255,255,255,0.7)', bordercolor='black', b
24     width=1000, height=600
25 )
26
27 fig.show()
28
```

Forecast vs Actuals





```

In [33]: 1 aligned_test = test.values[:len(fc)]
          2
          3 def forecast_accuracy(forecast, actual):
          4     mpe = np.mean((forecast - actual) / actual)
          5     rmse = np.sqrt(np.mean((forecast - actual) ** 2))
          6     minmax = 1 - np.mean(np.min([forecast, actual], axis=0) / np.max([forecast, actual]
          7
          8     return {'mpe': mpe, 'rmse': rmse, 'minmax': minmax}
          9
          10 accuracy_metrics = forecast_accuracy(fc, aligned_test)
          11 accuracy_metrics
          12

```

```

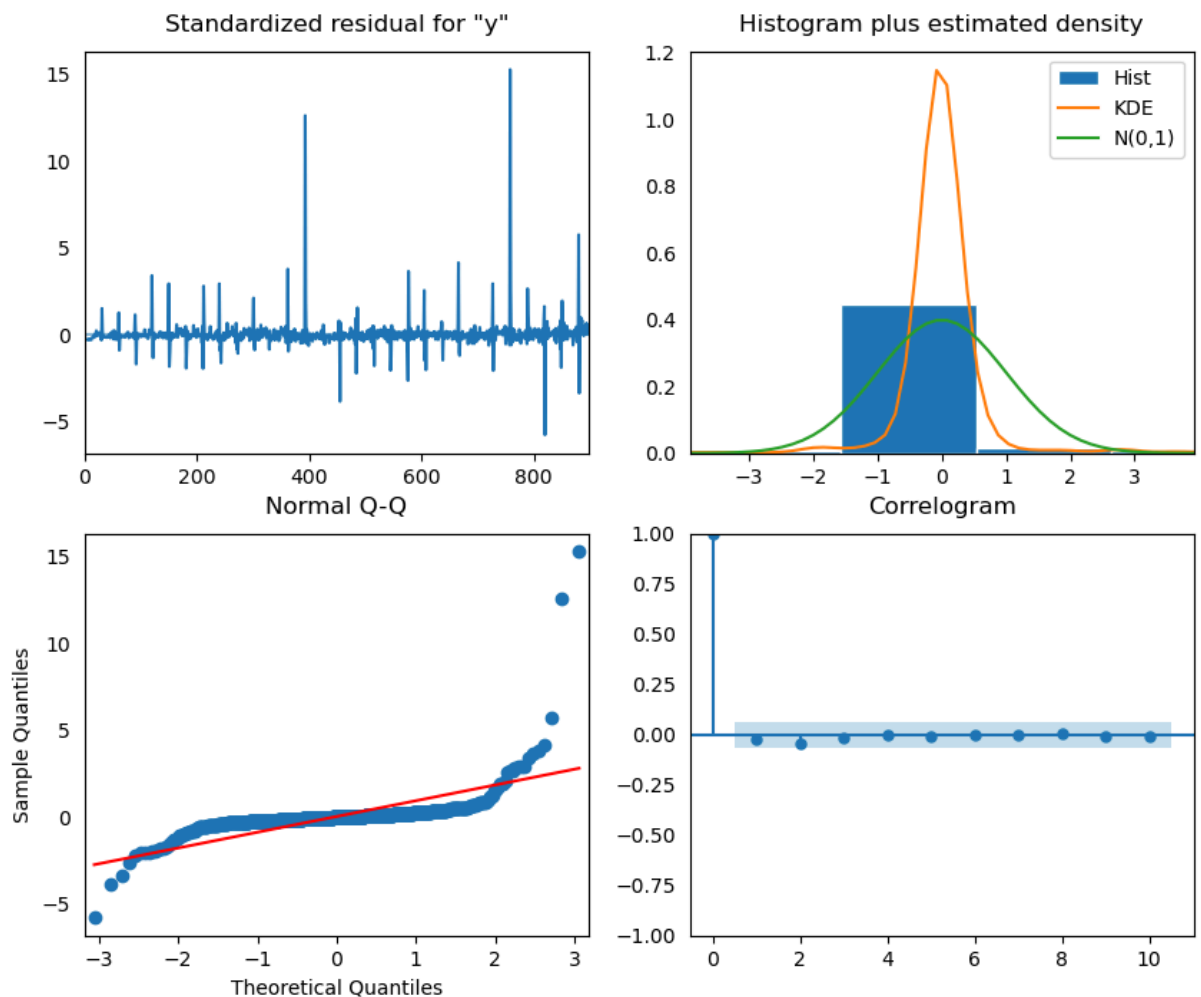
Out[33]: {'mpe': 1.1146247811868633,
          'rmse': 289658.2148717218,
          'minmax': 0.8344191602810813}

```

```

In [34]: 1 fitted.plot_diagnostics(figsize=(10,8))
          2 plt.show()

```



• **After training the model and removing seasonality from the data, we notice that**

- Standardized residual: The residual errors seem to fluctuate around a mean of zero and have a uniform - variance.
- Histogram: The density plot suggest normal distribution
- Theoretical Quantiles: Mostly the dots fall perfectly in line with the red line.
- Correlogram: The Correlogram, (or ACF plot) shows the residual errors are not autocorrelated.
- Overall, the model seems to be a good fit. So, let's use it to forecast

• **Final Forecasting**

```

In [35]: 1 steps = 150
2 forecast_result = fitted.get_forecast(steps=steps)
3
4 forecast = forecast_result.predicted_mean
5 conf_int = forecast_result.conf_int(alpha=0.05)
6
7 forecast_index = pd.date_range(start=test.index[-1], periods=steps + 1, freq='D')[1:]
8 forecast_series = pd.Series(forecast.values, index=forecast_index)
9
10 lower_series = pd.Series(conf_int.iloc[:, 0].values, index=forecast_index)
11 upper_series = pd.Series(conf_int.iloc[:, 1].values, index=forecast_index)
12
13
14 fig = go.Figure()
15
16 fig.add_trace(go.Scatter(x=train.index, y=train, mode='lines', name='Training', line=dict(
17 fig.add_trace(go.Scatter(x=test.index, y=test, mode='lines', name='test', line=dict(col
18
19 fig.add_trace(go.Scatter(x=forecast_series.index, y=forecast_series, mode='lines',
20 name='Forecast', line=dict(color='fuchsia'))))
21
22 fig.add_trace(go.Scatter(x=forecast_series.index, y=lower_series, mode='lines',
23 name='Lower Bound', line=dict(width=0), fill=None))
24 fig.add_trace(go.Scatter(x=forecast_series.index, y=upper_series, mode='lines',
25 name='Upper Bound', line=dict(width=0), fill='tonexty',
26 fillcolor='rgba(128, 128, 128, 0.2)', showlegend=False))
27
28 fig.update_layout(
29 title='Forecast vs Actuals with Future Predictions',
30 xaxis_title='Date',
31 yaxis_title='Total Due',
32 width=900, height=500,
33 legend=dict(yanchor='top', y=0.99, xanchor='left', x=0.01)
34 )
35
36 fig.show()
37

```

Forecast vs Actuals with Future Predictions

